

CSC311 A3, Winter 2019

Timothy Lee (leetim13)

1. a)

$$\begin{aligned}
 L(\theta) &= p(x, c | \theta, \pi) \\
 &= p(c | \theta, \pi) p(x | c, \theta, \pi) \\
 &= p(c | \pi) \prod_{j=1}^{784} p(x_j | c, \theta_{jc}) \\
 &= \pi_c \prod_{j=1}^{784} \theta_{jc}^{x_j} (1 - \theta_{jc})^{1-x_j}
 \end{aligned}$$

Taking logarithms for MLE,

$$\begin{aligned}
 l(\theta) &= \log(L(\theta)) \\
 &= \sum_{j=1}^N (\log \pi_j^{t_j^{(i)}} + \sum_{j=1}^{784} (x_j^{(i)} \log \theta_{jc} + (1 - x_j^{(i)}) \log(1 - \theta_{jc})))
 \end{aligned}$$

Calculating partial derivative to find *argmax*,

$$\frac{\partial l(\theta_{jc})}{\partial \theta_{jc}} = \sum_{i=1}^N \left(\frac{x_j^{(i)}}{\theta_{jc}} - \frac{1 - x_j^{(i)}}{1 - \theta_{jc}} \right)$$

Setting derivative to 0,

$$\sum_{i=1}^N \mathbb{I}(c^{(i)} = c) \theta_{jc} = \sum_{i=1}^N \mathbb{I}(c^{(i)} = c) x_j^{(i)}$$

Therefore,

$$\theta_{jc}^{MLE} = \frac{\sum_{i=1}^N \mathbb{I}(C^{(i)} = c) x_j^{(i)}}{\sum_{i=1}^N \mathbb{I}(C^{(i)} = c)}$$

Note that the indicator function $\mathbb{I}(c^{(i)} = c) x_j^{(i)} = \mathbb{I}(c^{(i)} = c \text{ and } x_j^{(i)} = 1)$. In other words, it is use to denote the number of times j -th pixel = 1 in the c -th class.

Now, we will find the MLE estimate for prior π .

First (by hint), we know that $L(\theta) = p(t^i|\pi) = \prod_{j=0}^9 \pi_j^{t_j^{(i)}}$, where $\sum_{j=0}^9 \pi_j = 1$. Applying logarithms,

$$l(\theta) = t_j^{(i)} \sum_{j=0}^9 \log(\pi_j)$$

Since (by hint) we can denote $\pi_9 = 1 - \sum_{i=0}^8 \pi_i$. We can re-write $l(\theta)$ as

$$l(\theta) = \sum_{i=0}^N (t_9^{(i)} \log(1 - \sum_{j=0}^8 \pi_j) + \sum_{j=0}^8 t_j^{(i)} \log(\pi_j))$$

Taking derivative with respect to π_j :

$$\frac{\partial l(\theta)}{\partial \pi} = \sum_{i=0}^N (t_j^{(i)} \frac{1}{\pi_j} + t_9^{(i)} - \frac{1}{1 - \sum_{j=0}^8 \pi_j}), \text{ for } j = 0, \dots, 8$$

Setting derivative to 0, we have

$$\sum_{i=0}^N \frac{t_j^{(i)}}{\pi_j} = \sum_{i=0}^N \frac{t_9^{(i)}}{\pi_9}$$

Then,

$$\begin{aligned} \frac{\hat{\pi}_j}{\hat{\pi}_9} &= \frac{\sum_{i=0}^N t_j^{(i)}}{\sum_{i=0}^N t_9^{(i)}} \\ \hat{\pi}_j &= \frac{\sum_{i=0}^N t_j^{(i)}}{\sum_{i=0}^N t_9^{(i)}} \cdot \hat{\pi}_9 \end{aligned}$$

Since $\pi_9 = 1 - \sum_{j=0}^8 \pi_j$ (by hint) and $\sum_{j=0}^9 t_j^{(i)} = 1$ for each class label $t^{(i)}$ (by definition of 1-of-10 encoded class), we can conclude that

$$\begin{aligned} \hat{\pi}_9 &= 1 - \sum_{j=0}^8 \hat{\pi}_j \\ &= 1 - \sum_{j=0}^8 \frac{\sum_{i=0}^N t_j^{(i)}}{\sum_{i=0}^N t_9^{(i)}} \cdot \hat{\pi}_9 \quad (\text{by definition of } \hat{\pi}_j \text{ above}) \\ &= 1 - \frac{\sum_{i=0}^N (1 - t_9^{(i)})}{\sum_{i=0}^N t_9^{(i)}} \cdot \hat{\pi}_9 \quad (\text{since } \sum_{j=0}^9 t_j^{(i)} = 1) \\ &= \frac{\sum_{i=0}^N t_9^{(i)}}{N} \quad (\text{by re-arranging}) \end{aligned}$$

Therefore,

$$\begin{aligned}\pi_j^{MLE} &= \frac{\sum_{i=0}^N t_j^{(i)}}{\sum_{i=0}^N t_9^{(i)}} \cdot \frac{\sum_{i=0}^N t_9^{(i)}}{N} \quad (\text{by substitution}) \\ &= \frac{\sum_{i=0}^N t_j^{(i)}}{N} \\ &= \frac{\sum_{i=0}^N \mathbb{I}(t_c^{(i)} = t_j)}{N}\end{aligned}$$

Again, the indicator function is used to denote the number of times j -th pixel = 1 in the c -th class and N is the total number of training samples.

b) We can denote $L(\theta)$ as

$$p(c|x, \theta, \pi) = \frac{p(c|\pi)p(x|c, \theta)}{p(x|\theta, \pi)}$$

Then, the log-likelihood, $l(\theta)$ would be

$$\begin{aligned}\log p(c|x, \theta, \pi) &= \log(p(c|\pi)) + \log(p(x|c, \theta)) - \log(p(x|\theta, \pi)) \\ &= \log(p(c|\pi)) + \log\left(\prod_{j=1}^{784} p(x_j|c, \theta_{jc})\right) + \log\left(\sum_{j=0}^9 p(x|c = j, \theta, \pi)p(c = j)\right)\end{aligned}$$

Hence, $\log p(t|x, \theta, \pi)$ can be written as

$$= \sum_{j=0}^9 t_j^{(i)} \log \pi_j + \sum_{j=1}^{784} (x_j \log \theta_{jc} + (1 - x_j) \log(1 - \theta_{jc})) - \log\left(\sum_{j=0}^9 \pi_j \prod_{j=1}^{784} \theta_{jc}^{x_j} (1 - \theta_{jc})^{(1-x_j)}\right)$$

c) We encounter a divide by 0 error when performing logarithm.

```
C:/Users/hwtim/Desktop/311a3/naive_bayes.py:129: RuntimeWarning: divide by zero encountered in log
print(np.log(theta).shape)
```

Figure 1: Error Message

d)




Figure 2: MLE estimator

e) For deriving the *Maximum A posterior Proability (MAP)* given $\theta_{jc} \sim \beta(3, 3)$, we will first derive the pdf of this prior beta distribution,

$$f(\theta_{jc}) = \frac{1}{B(3, 3)} \theta_{jc}^{(3-1)} (1 - \theta_{jc})^{(3-1)}$$

Ignoring beta normalizing constant (as per question) and simplifying,

$$f(\theta_{jc}) = \theta_{jc}^2 (1 - \theta_{jc})^2$$

Now applying Baye's Theorem,

$$\begin{aligned} p(\theta_{jc} | x_j, c, \pi) &\propto p(\theta_{jc}) p(x_j, c | \theta_{jc}, \pi) \\ &\propto \theta_{jc}^2 (1 - \theta_{jc})^2 \pi_c^i \prod_{i=1}^{784} \theta_{jc}^{x_{ij}} (1 - \theta_{jc})^{(1-x_{ij})} \end{aligned}$$

Taking logarithms,

$$l(\theta) = 2 \log(\theta_{jc}) + 2 \log(1 - \theta_{jc}) + \sum_{i=1}^{784} x_{ij} \log \theta_{jc} + (1 - \theta_{jc}) \log (1 - x_{ij}) + k$$

, where k is a constant factor.

Again, I will use $\mathbb{I}(C^i = c)$ to denote $x_j = 1$ in the c -th class.

Taking derivatives to find *argmax* of θ ,

$$\frac{\partial l}{\partial \theta_{jc}} = \left(\frac{2}{\theta_{jc}} - \frac{2}{1 - \theta_{jc}} \right) + \sum_{i=1}^N \mathbb{I}(C^i = c) \left(\frac{x_j^{(i)}}{\theta_{jc}} - \frac{1 - x_j^{(i)}}{1 - \theta_{jc}} \right)$$

Setting derivative to 0,

$$-2\theta_{jc} + (2 - 2\theta_{jc}) + \sum_{i=1}^N \mathbb{I}(C^i = c) (x_j^{(i)} (1 - \theta_{jc}) - (1 - x_j^{(i)}) \theta_{jc}) = 0$$

Re-arranging,

$$-4\theta_{jc} + 2 + \sum_{i=1}^N \mathbb{I}(C^i = c)(x_j^{(i)} - x_j^{(i)}\theta_{jc} - \theta_{jc} + x_j^{(i)}\theta_{jc}) = 0$$

Finally, we get,

$$\theta_{jc}^{MAP} = \frac{\sum_{i=1}^N \mathbb{I}(C^i = c)x_j^{(i)} + 2}{\sum_{i=1}^N \mathbb{I}(C^i = c) + 4}$$

f)

```
Average log-likelihood for MLE is nan
Average log-likelihood for MAP is -3.3570631378602687
Training accuracy for MAP is 0.8352166666666667
Test accuracy for MAP is 0.816
```

Figure 3: Accuracy and average log-likelihood

g)




Figure 4: Theta MAP

```

1 from __future__ import absolute_import
2 from __future__ import print_function
3 from future.standard_library import install_aliases
4 install_aliases()
5 import numpy as np
6 import os
7 import gzip
8 import struct
9 import array
10 import matplotlib.pyplot as plt
11 import matplotlib.image
12 from urllib.request import urlretrieve
13 from scipy.special import logsumexp
14
15 def download(url, filename):
16     if not os.path.exists('data'):
17         os.makedirs('data')
18     out_file = os.path.join('data', filename)
19     if not os.path.isfile(out_file):
20         urlretrieve(url, out_file)
21
22 def mnist():
23     base_url = 'http://yann.lecun.com/exdb/mnist/'
24
25     def parse_labels(filename):
26         with gzip.open(filename, 'rb') as fh:
27             magic, num_data = struct.unpack(">II", fh.read(8))
28             return np.array(array.array("B", fh.read()), dtype=np.uint8)
29
30     def parse_images(filename):
31         with gzip.open(filename, 'rb') as fh:
32             magic, num_data, rows, cols = struct.unpack(">IIII", fh.read(16))
33             return np.array(array.array("B", fh.read()), dtype=np.uint8).reshape(num_data, rows, cols)
34
35     for filename in ['train-images-idx3-ubyte.gz',
36                     'train-labels-idx1-ubyte.gz',
37                     't10k-images-idx3-ubyte.gz',
38                     't10k-labels-idx1-ubyte.gz']:
39         download(base_url + filename, filename)
40
41     train_images = parse_images('data/train-images-idx3-ubyte.gz')
42     train_labels = parse_labels('data/train-labels-idx1-ubyte.gz')
43     test_images = parse_images('data/t10k-images-idx3-ubyte.gz')
44     test_labels = parse_labels('data/t10k-labels-idx1-ubyte.gz')
45     return train_images, train_labels, test_images[:1000], test_labels[:1000]
46
47
48 def load_mnist():
49     partial_flatten = lambda x: np.reshape(x, (x.shape[0], np.prod(x.shape[1:])))
50     one_hot = lambda x, k: np.array(x[:, None] == np.arange(k)[None, :], dtype=int)
51     train_images, train_labels, test_images, test_labels = mnist()
52     train_images = (partial_flatten(train_images) / 255.0 > .5).astype(float)
53     test_images = (partial_flatten(test_images) / 255.0 > .5).astype(float)
54     train_labels = one_hot(train_labels, 10)
55     test_labels = one_hot(test_labels, 10)
56     N_data = train_images.shape[0]
57     return N_data, train_images, train_labels, test_images, test_labels

```

Figure 5: Code Part 1

```

60 def plot_images(images, ax, ims_per_row=5, padding=5, digit_dimensions=(28, 28),
61                 cmap=matplotlib.cm.binary, vmin=None, vmax=None):
62     """Images should be a (N_images x pixels) matrix."""
63     N_images = images.shape[0]
64     N_rows = np.int32(np.ceil(float(N_images) / ims_per_row))
65     pad_value = np.min(images.ravel())
66     concat_images = np.full(((digit_dimensions[0] + padding) * N_rows + padding,
67                               (digit_dimensions[1] + padding) * ims_per_row + padding), pad_value)
68     for i in range(N_images):
69         cur_image = np.reshape(images[i, :], digit_dimensions)
70         row_ix = i // ims_per_row
71         col_ix = i % ims_per_row
72         row_start = padding + (padding + digit_dimensions[0]) * row_ix
73         col_start = padding + (padding + digit_dimensions[1]) * col_ix
74         concat_images[row_start: row_start + digit_dimensions[0],
75                       col_start: col_start + digit_dimensions[1]] = cur_image
76         cax = ax.matshow(concat_images, cmap=cmap, vmin=vmin, vmax=vmax)
77         plt.xticks(np.array([]))
78         plt.yticks(np.array([]))
79     return cax
80
81
82 def save_images(images, filename, **kwargs):
83     fig = plt.figure(1)
84     fig.clf()
85     ax = fig.add_subplot(111)
86     plot_images(images, ax, **kwargs)
87     fig.patch.set_visible(False)
88     ax.patch.set_visible(False)
89     plt.savefig(filename)
90
91
92 def train_mle_estimator(train_images, train_labels):
93     """ Inputs: train_images, train_labels
94         Returns the MLE estimators theta_mle and pi_mle"""
95     theta = np.matmul(np.transpose(train_images), train_labels) # 784 x 10 array
96     sum_labels = train_labels.sum(axis=0) #sum of indicator function
97     theta_mle = theta/sum_labels
98
99     N = len(train_images)
100     pi_mle = np.ones(N)
101     pi_mle = pi_mle.dot(train_labels) * 1/N
102     return theta_mle, pi_mle
103
104
105 def train_map_estimator(train_images, train_labels):
106     """ Inputs: train_images, train_labels
107         Returns the MAP estimators theta_map and pi_map"""
108
109     # YOU NEED TO WRITE THIS PART
110     N = len(train_images)
111     theta = np.matmul(np.transpose(train_images), train_labels) # 784 x 10 array
112     sum_labels = train_labels.sum(axis=0) #sum of indicator function
113     pi_map = sum_labels/N
114     theta_map = (theta + 2)/(4 + sum_labels) #using beta_prior
115     return theta_map, pi_map

```

Figure 6: Code Part 2

```

118 def log_likelihood(images, theta, pi):
119     """ Inputs: images, theta, pi
120     Returns the matrix 'log_like' of loglikelihoods over the input images where
121     log_like[i,c] = log p(c|x^(i), theta, pi) using the estimators theta and pi.
122     log_like is a matrix of num of images x num of classes
123     Note that log likelihood is not only for c^(i), it is for all possible c's."""
124     N_pi = len(pi)
125     N = len(images)
126     shape = (N, N_pi)
127     log_like = np.zeros(shape)
128     log_p_x = logsumexp(np.log(pi) + np.dot(images, np.log(theta)) + np.dot((1. - images), np.log(1. - theta)), axis=1)
129     for c in range(N_pi):
130         log_like[:, c] = (np.dot(images, np.log((theta.T)[c])) + np.dot((1. - images), np.log(1. - (theta.T)[c]))
131             - log_p_x) + np.log(pi[c]) #by derivation
132     return log_like
133
134
135 def predict(log_like):
136     """ Inputs: matrix of log likelihoods
137     Returns the predictions based on log likelihood values"""
138
139     # YOU NEED TO WRITE THIS PART
140     predictions = np.argmax(log_like, axis =1)
141     return predictions
142
143
144 def accuracy(log_like, labels):
145     """ Inputs: matrix of log likelihoods and 1-of-K labels
146     Returns the accuracy based on predictions from log likelihood values"""
147     N = len(labels)
148     one_of_k = labels[np.arange(N),log_like.argmax(1)]
149     accuracy = sum(one_of_k)/N
150     return accuracy
151
152
153 def image_sampler(theta, pi, num_images):
154     """ Inputs: parameters theta and pi, and number of images to sample
155     Returns the sampled images"""
156
157     # YOU NEED TO WRITE THIS PART
158     l = len(theta)
159     shape = num_images, l
160     sampled_images = np.zeros(shape) #new array of given shape and type, filled with zeros.
161     for i in range(num_images):
162         c = np.random.choice(10, p=pi) #by hint
163         sampled_images[i] = np.random.binomial(1, p=theta[:,c]) #by hint
164         save_images(sampled_images, "samples.png")
165     return sampled_images
166

```

Figure 7: Code Part 3


```

168 def main():
169     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
170
171     # Fit MLE and MAP estimators
172     theta_mle, pi_mle = train_mle_estimator(train_images, train_labels)
173     theta_map, pi_map = train_map_estimator(train_images, train_labels)
174
175     # Find the log likelihood of each data point
176     loglike_train_mle = log_likelihood(train_images, theta_mle, pi_mle)
177     loglike_train_map = log_likelihood(train_images, theta_map, pi_map)
178
179
180     avg_loglike_mle = np.sum(loglike_train_mle * train_labels) / N_data
181     avg_loglike_map = np.sum(loglike_train_map * train_labels) / N_data
182 #
183     print("Average log-likelihood for MLE is ", avg_loglike_mle)
184     print("Average log-likelihood for MAP is ", avg_loglike_map)
185 #
186     train_accuracy_map = accuracy(loglike_train_map, train_labels)
187     loglike_test_map = log_likelihood(test_images, theta_map, pi_map)
188     test_accuracy_map = accuracy(loglike_test_map, test_labels)
189 #
190     print("Training accuracy for MAP is ", train_accuracy_map)
191     print("Test accuracy for MAP is ", test_accuracy_map)
192
193     # Plot MLE and MAP estimators
194 #     print(theta_mle.T.shape)
195     save_images(theta_mle.T, 'mle.png')
196     save_images(theta_map.T, 'map.png')
197
198     # Sample 10 images
199     sampled_images = image_sampler(theta_map, pi_map, 10)
200     save_images(sampled_images, 'sampled_images.png')
201
202
203 if __name__ == '__main__':
204     main()
205

```

Figure 8: Code Part 4

2. a) (Trivially) True, under the given assumptions of the naive bayes model where we assume the independence of x_i, x_j , which might not be the case in the real world.

b) False, we will show that x_i and x_j are not independent by showing $p(x_i, x_j) \neq p(x_i)p(x_j)$
When marginalizing over c , we have

$$p(x_i, x_j) = \sum p(x_i, x_j | c) = \sum p(x_i | c) p(x_j | c)$$

, while

$$p(x_i)p(x_j) = \sum p(x_i | c) \sum p(x_j | c)$$

c)




Figure 9: Random Image Samples

```

166 def image_sampler(theta, pi, num_images):
167     """ Inputs: parameters theta and pi, and number of images to sample
168     Returns the sampled images"""
169
170     # YOU NEED TO WRITE THIS PART
171     l = len(theta)
172     shape = num_images, l
173     sampled_images = np.zeros(shape) #new array of given shape and type, filled with zeros.
174     for i in range(num_images):
175         c = np.random.choice(10, p=pi) #by hint
176         sampled_images[i] = np.random.binomial(1, p=theta[:,c]) #by hint
177     save_images(sampled_images, "samples.png")
178     return sampled_images
179

```

Figure 10: Code for image sampler

- First, I will provide some intuition of PCA in terms of graphs and images (as described, but not required by the question).




Figure 11: First 3 PCA of digit '2'




Figure 12: First 3 PCA of digit '3'

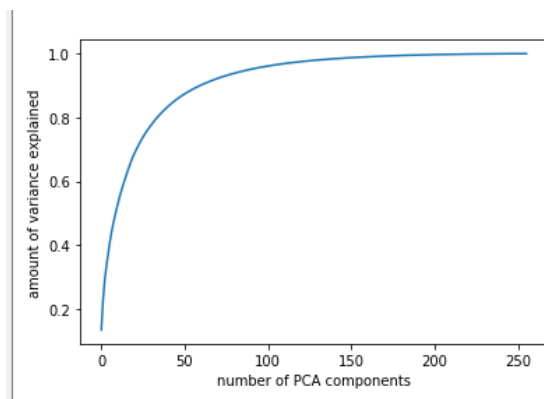


Figure 13: Graph of number of PCA components vs variance explained

a)




Figure 14: Graph of validation set classification error rates versus number of eigenvectors, K

b) Based on the graph, we could conclude that $K = 20$ seems like a reasonable number of eigen vector to choose from, since the classification error rate on the validation set is the lowest (global minimum), implying a higher classification accuracy. We can deduce that as K continues to increase from 2, the classification error reduces up until $K = 20$, where the error appears to go up again. This is most likely due to overfitting of our training data (which causes the performance of the training data set to be almost perfect, but the performance to be a lot worse in our validation data).

c) As argued before, $K = 20$ seems like a logical choice, which has about an error rate of around 1%. This is the optimum/lowest error rate achieved in the validation set for K in which we avoid over-fitting.

Error of $K=20 = 0.010000000000000009$

Figure 15: Error for $K=20$

```

1 import numpy as np
2 from sklearn.decomposition import PCA #used to build intuition
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5
6 def load_data(filename, load2=True, load3=True):
7     """Loads data for 2's and 3's
8     Inputs:
9         filename: Name of the file.
10        load2: If True, load data for 2's.
11        load3: If True, load data for 3's.
12    """
13    assert (load2 or load3), "Atleast one dataset must be loaded."
14    data = np.load(filename)
15    if load2 and load3:
16        inputs_train = np.hstack((data['train2'], data['train3']))
17        inputs_valid = np.hstack((data['valid2'], data['valid3']))
18        inputs_test = np.hstack((data['test2'], data['test3']))
19        target_train = np.hstack((np.zeros((1, data['train2'].shape[1])), np.ones((1, data['train3'].shape[1]))))
20        target_valid = np.hstack((np.zeros((1, data['valid2'].shape[1])), np.ones((1, data['valid3'].shape[1]))))
21        target_test = np.hstack((np.zeros((1, data['test2'].shape[1])), np.ones((1, data['test3'].shape[1]))))
22    else:
23        if load2:
24            inputs_train = data['train2']
25            target_train = np.zeros((1, data['train2'].shape[1]))
26            inputs_valid = data['valid2']
27            target_valid = np.zeros((1, data['valid2'].shape[1]))
28            inputs_test = data['test2']
29            target_test = np.zeros((1, data['test2'].shape[1]))
30        else:
31            inputs_train = data['train3']
32            target_train = np.zeros((1, data['train3'].shape[1]))
33            inputs_valid = data['valid3']
34            target_valid = np.zeros((1, data['valid3'].shape[1]))
35            inputs_test = data['test3']
36            target_test = np.zeros((1, data['test3'].shape[1]))
37
38    return inputs_train.T, inputs_valid.T, inputs_test.T, target_train.T, target_valid.T, target_test.T
39
40 #Intution of displaying PCA compoenents vs images
41 def show(g, imshape, i, j, x, title=None):
42     ax = fig.add_subplot(g[i, j], xticks=[], yticks=[])
43     ax.imshow(x.reshape(imshape), interpolation='nearest')
44     if title:
45         ax.set_title(title, fontsize=12)
46

```

Figure 16: Code

```

47 def plot_pca_components(x, coefficients=None, mean=0, components=None,
48                         imshape=(16, 16), n_components=8, fontsize=12,
49                         show_mean=True):
50     '''Building intuition by viewing top k PCA
51     ...
52     if coefficients is None:
53         coefficients = x
54     if components is None:
55         components = np.eye(len(coefficients), len(x))
56     mean = np.zeros_like(x) + mean
57     fig = plt.figure(figsize=(1.2 * (5 + n_components), 1.2 * 2))
58     g = plt.GridSpec(2, 4 + bool(show_mean) + n_components, hspace=0.3)
59     show(g, imshape, slice(2), slice(2), x, "True Original Image")
60     approx = mean.copy()
61     counter = 2
62     if show_mean:
63         show(g, imshape, 0, 2, np.zeros_like(x) + mean, r'$mean$')
64         show(g, imshape, 1, 2, approx)
65         counter += 1
66     for i in range(n_components):
67         approx = approx + coefficients[i] * components[i]
68         show(g, imshape, 0, i + counter, components[i], r'$k_{0}$'.format(i + 1))
69         show(g, imshape, 1, i + counter, approx,
70              r"${0:.2f} \cdot c_{1}$".format(coefficients[i], i + 1))
71         if show_mean or i > 0:
72             plt.gca().text(0, 1.05, '$+$', ha='right', va='bottom',
73                            transform=plt.gca().transAxes, fontsize=fontsize)
74     show(g, imshape, slice(2), slice(-2, None), approx, "Approx")
75     return fig
76
77 def plot_digits(data):
78     fig, axes = plt.subplots(10, 10, figsize=(10, 4),
79                             subplot_kw={'xticks':[], 'yticks':[]},
80                             gridspec_kw=dict(hspace=0.1, wspace=0.1))
81     for i, ax in enumerate(axes.flat):
82         ax.imshow(data[i].reshape(16, 16), #reshape into 16x16 in order to be displayed
83                  cmap='binary', interpolation='nearest')
84

```

Figure 17: Code

```

84
85 def first_k_components(training_data, k):
86     ...
87     Plot of first k_components vs eigen values
88     ...
89     mean = np.mean(training_data, axis =0)
90     num_repeated = training_data.shape[0], 1
91     centered_data = training_data - np.tile(mean, num_repeated) #subtract the mean of training data
92     data_T = centered_data.T
93     covariance_matrix = np.cov(data_T)
94     eigen_values_cov , eigen_vectors_cov = np.linalg.eig(covariance_matrix )
95     eigen_values_cov = eigen_values_cov[:, -1]
96     eigen_vectors_cov = eigen_vectors_cov[:, -1]
97     plt.figure()
98     length = np.arange(0.0 , len(eigen_values_cov ) , 1) #[0,1,...,256]
99     plt.plot(length, eigen_values_cov )
100    plt.xlabel ("number of eigenvectors")
101    plt.ylabel ("accuracy")
102    plt.title ("plot of eigenvalues of covariance")
103    plt.grid ( True )
104    eigen_values = eigen_values_cov[:k]
105    eigen_vectors = eigen_vectors_cov[:k ,:]
106    # print(eigen_values.shape)  #(10,)
107    # print(eigen_vectors.shape)  #(10, 256)
108    return eigen_values, eigen_vectors , mean
109
110 def one_nn_classifier(train_data, train_labels, valid_data, k=1) :
111     N = len(valid_data)
112     shape = N, 1
113     valid_labels = np.zeros(shape) #initialize empty
114     train_data_N = len(train_data)
115
116     for i in range (N):
117         min_index = -1
118         min_value = np.inf
119         for j in range (train_data_N):
120             euclidean_distance = np.linalg.norm(valid_data[i]- train_data[j])
121             if euclidean_distance < min_value :
122                 min_value = euclidean_distance
123                 min_index = j
124             valid_labels[i] = train_labels[min_index]
125     return valid_labels
126

```

Figure 18: Code

```

126
127 def extract_eigen_features(training_data):
128     mean = np.mean(training_data , axis =0)
129     num_repeated = training_data.shape[0], 1
130     centered_data = training_data - np.tile(mean, num_repeated)
131     data_T = centered_data.T
132     covariance_matrix = np.cov(data_T)
133     eigen_values_cov , eigen_vectors_cov = np.linalg.eig(covariance_matrix )
134     # print(eigen_values_cov.shape) #(10,)
135     # print(eigen_vectors_cov.shape) #(10, 256)
136     sorted_eigen_values = eigen_values_cov.argsort()[::-1] #vector of sorted eigen values asc
137     eigen_values = eigen_values_cov[sorted_eigen_values]
138     eigen_vectors = eigen_vectors_cov[:,sorted_eigen_values]
139     return eigen_values , eigen_vectors , mean
140
141 def accuracy (prediction_value, target_value):
142     return np.mean(target_value==prediction_value)
143
144 def train_model_pca(given_K , inputs_train , inputs_valid , target_train , target_valid):
145     accuracy_list = []
146     eigen_values , eigen_vectors , mean = extract_eigen_features(inputs_train)
147     for k in given_K :
148         code_vectors = eigen_vectors[:, : k]
149         # print(top_k_vector)
150         top_k_value = value[:, : k]
151         num_repeated_training = (inputs_train.shape[0] , 1)
152         # print(num_repeated_training.shape) #600
153         centered_training_data = inputs_train - np.tile(mean, num_repeated_training )
154         num_repeated_valid = (inputs_valid.shape[0] , 1)
155         # print(num_repeated_valid.shape) #200
156         centered_valid_data = inputs_valid - np .tile(mean, num_repeated_valid)
157
158         #projection onto the low-dimensional space
159         low_dim_space_valid = np.dot(centered_valid_data, code_vectors)
160         low_dim_space_train = np.dot(centered_training_data, code_vectors)
161
162         #using 1-NN classifier on K dimensional features
163         low_dim_space_target = one_nn_classifier(low_dim_space_train , target_train , low_dim_space_valid )
164         accuracy_ = accuracy(low_dim_space_target, target_valid)
165         error = 1 - accuracy_
166         accuracy_list.append(error)
167
168 plt.figure()
169 plt.grid(True)
170 plt.plot(given_K , accuracy_list)
171 plt.xlabel('first K principal components')
172 plt.ylabel('Classification Error rates')
173 plt.title ('plot of accuracy vs. #of eigen vectors')
174 return accuracy_list

```

Figure 19: Code


```

176 if __name__ == '__main__':
177     inputs_train, inputs_valid, inputs_test, target_train, target_valid, target_test = load_data("digits.npz")
178     # print(inputs_train.shape)
179     # print(inputs_valid.shape)
180     # print(target_train.shape)
181     # print(target_valid.shape)
182
183     given_K = [2, 5, 10, 20, 30]
184     # view_eig_vector_images(10, top_k_vector, mean)
185     accuracy_list = train_model_pca(given_K, inputs_train, inputs_valid, target_train, target_valid)
186
187     # print(accuracy_k)
188     best_K = 20 #after selection
189     accuracy_list = train_model_pca([best_K], inputs_train, inputs_test, target_train, target_test)
190     print("Error of K=" + str(best_K) + " = " + str(accuracy_list[0]))
191
192     pca = PCA().fit(inputs_train) #only used to build intuition
193     plt.plot(np.cumsum(pca.explained_variance_ratio_))
194     plt.xlabel('number of PCA components')
195     plt.ylabel('amount of variance explained')
196     pca = PCA(n_components=20)
197     Xproj = pca.fit_transform(inputs_train)
198     fig = plot_pca_components(inputs_train[155], Xproj[155],
199                             pca.mean_, pca.components_)
200     plot_digits(inputs_train)
201     plt.show
202

```

Figure 20: Code